Z3str3: A String Solver with Theory-aware Branching

Murphy Berzish[1], Yunhui Zheng[2], and Vijay Ganesh[1]

[1] University of Waterloo, Waterloo, Canada
[2] IBM Research, Yorktown Heights, USA

Abstract. We present a new string SMT solver, Z3str3, that is faster than its competitors Z3str2, Norn, CVC4, S3, and S3P over majority of three industrial-strength benchmarks, namely, Kaluza, PISA, and IBM AppScan. Z3str3 supports string equations, linear arithmetic over length function, and regular language membership predicate. The key algorithmic innovation behind the efficiency of Z3str3 is a technique we call theory-aware branching, wherein we modify Z3's branching heuristic to take into account the structure of theory literals to compute branching activities. In the traditional DPLL(T) architecture, the structure of theory literals is hidden from the DPLL(T) SAT solver because of the Boolean abstraction constructed over the input theory formula. By contrast, the theory-aware technique presented in this paper exposes the structure of theory literals to the DPLL(T) SAT solver's branching heuristic, thus enabling it to make much smarter decisions during its search than otherwise. As a consequence, Z3str3 has better performance than its competitors.

1   Introduction

String SMT solvers are increasingly becoming important for security applications and in the context of analysis of string-intensive programs [4,6,7,10,11,14,18]. Many string SMT solvers, such as Z3str2 [19,20] (and its predecessor Z3str [21]), CVC4 [8], Norn [2], S3 [16] (and its successor S3P [17]), and Stranger (and its successor ABC [3]) have been developed to address these challenges and applications. We have developed the Z3str3 string solver as a native first-class theory solver directly integrated into the Z3 SMT solver [5], that is much faster than its predecessors Z3str2 and Z3str, as well as competitors CVC4, Norn, and S3. Having direct access to the core solver of Z3 has allowed us to develop and implement a novel DPLL(T) technique which we call *theory-aware branching*, described below. We follow the latest string SMT language standard supported by all major string solvers, and published on the CVC4 website [8].

1.1   Contributions

1. **Theory-aware branching:** We leverage the integration between the Z3 SMT solver's DPLL(T) SAT layer (henceforth referred to as the core solver)

and the string solver to guide the search and prioritize certain branches of the search tree over others. In particular, we modify the activity computations of the branching heuristic of the Z3 core solver, making the heuristic aware of the structure of the theory literals underlying the Boolean abstraction of the input formula such that "simpler" theory literals are prioritized over more complex ones.

2. **Theory-aware case-split:** We add an optimization to Z3's core solver that enables efficient representation of mutually exclusive Boolean variables in the Boolean abstraction of the input theory formula.

3. **Experimental evaluation:** To validate the effectiveness of our techniques, we present a comprehensive and thorough evaluation of Z3str3, and compare against Z3str2, CVC4, S3, and Norn on several large industrial-strength benchmarks. We couldn't directly compare against S3P since its source is not available, but do summarize the results from their CAV 2016 paper and compare against Z3str3. We also couldn't compare against Stranger/ABC because these tools don't produce models for SAT cases, don't support dis-equations over arbitrary string terms, and have incorrectness issues as noted in their own paper [3].

2 Theory-Aware Branching

Several of the key enhancements we make in Z3str3 over Z3str2 involve changes to the core DPLL(T) SAT solver in Z3, which handles the Boolean structure of the formula and performs propagation and branching. The first of these enhancements is referred to as **theory-aware branching**. We modify the Z3 core solver to allow theory solvers to provide information about certain theory literals that are given increased or decreased priority during the search. For example, consider the case where the solver learns the equality $X \cdot Y = A \cdot B$ for non-constant terms X, Y, A, B. The behaviour of Z3str3 (in line with Z3str2) is to handle this equality by considering a disjunction of three possible arrangements [19,20]:

- $X = A$ and $Y = B$
- $X = A \cdot s_1$ and $s_1 \cdot Y = B$ for a fresh non-empty string variable s_1
- $X \cdot s_2 = A$ and $Y = s_2 \cdot B$ for a fresh non-empty string variable s_2

Of the three possible arrangements, the first is the simplest to check because it does not introduce any new variables and only asserts equalities between existing terms. Therefore, we would like Z3's core solver to prioritize checking this arrangement before the others. The advantage gained by theory-aware branching is the ability to give the core solver information regarding the relative importance of each branch, allowing the theory solver to exert additional control over the search. Note we always prioritize simpler branches over more complex ones.

We implement theory-aware branching as a modification of the branching heuristic in Z3. This idea of creating a theory-aware DPLL branching heuristic is mentioned in [13]. The default branching heuristic in Z3 is activity-based, similar to VSIDS [9]. The core solver will branch on the literal with the highest activity

that has not yet been assigned. In the VSIDS branching heuristic, activity is increased additively when a literal appears in a conflict clause, and decayed multiplicatively at regular intervals.

The theory-aware branching technique computes the activity of a literal A as the sum of two terms A_b and A_t, wherein the term A_b is the "base activity", which is the standard activity of the literal as computed and handled by Z3's core solver. The term A_t is the "theory-aware activity". The value of this term is provided for individual literals by theory solvers, and is taken to be 0 if no theory-aware branching information has been provided. This modification causes the core solver to branch on the literal with the highest activity A, taking into account both the standard activity value and the theory-aware activity. Therefore, assigning a (small) positive theory-aware activity to a literal will cause it to have higher activity than usual, making it more likely for the core solver to choose it to branch on. Conversely, assigning a (small) negative theory-aware activity will deter the core solver from choosing that literal. Theory-aware branching in Z3str3 modifies the activities of theory literals as follows:

1. Literals corresponding to arrangements that do not create new variables are given a large (0.5) A_t. Other arrangements are given a small (0.1) A_t.
2. Arrangements that allow a variable to become equal to a constant string are given a small (0.2) A_t.
3. When searching for length of strings, we give the literal corresponding to the choice "generate more length options" a small negative (-0.1) A_t.

3 Theory-Aware Case-split

During the search, a theory solver can create terms which encode a disjunction of Boolean literals that are pairwise mutually exclusive, i.e. exactly one of the literals must be assigned true and the others must be assigned false. We refer to this as a **theory-aware case-split**. As an example, consider the case where the string solver learns that a concatenation of two arbitrary string terms X and Y is equal to a string constant $c = c_1 c_2 \ldots c_n$ of length n, where each c_i is a character in c. At this point there are $n + 1$ possible ways in which we can split the constant c over X and Y resulting in different arrangements:

- $X = \epsilon, Y = c_1 c_2 \ldots c_n$
- $X = c_1, Y = c_2 c_3 \ldots c_n$
- $\ldots$
- $X = c_1 c_2 \ldots c_n, Y = \epsilon$

Note that each of these arrangements represents a case that can be explored by the solver, and also that all of these cases are mutually exclusive (as clearly X cannot be equal to both ϵ and c_1 simultaneously, etc.). Thus, this represents a theory-aware case-split. Note that the Boolean abstraction constructed over theory literals completely hides this obviously useful information that these

variables (and corresponding arrangements) are mutually exclusive. A naive solution is to encode $O(n^2)$ extra mutual exclusion Boolean clauses over these variables. Unfortunately, this would result in very poor performance because of the quadratic blowup in formula size. Another option is to let the congruence closure solver in the Z3 core discover the mutual exclusivity of these Boolean variables. This can result in unnecessary backtracking, unnecessary calls to congruence closure, and, in the worst case, reduces to the same set of mutual exclusion clauses being learned in the form of conflict clauses. We improve the performance of theory case-splits by allowing theory solvers to provide extra information to the core solver regarding which literals can be treated as mutually exclusive during its search. This means that theory solvers do not have to assert extra clauses to enforce mutual exclusivity of choices. Instead, our modified core solver can now use this extra information from the theory solver during its search. Our implementation of this technique is as follows:

1. The theory solver provides the core solver with a set S of mutually exclusive literals that correspond to a theory case-split. This set is maintained by the core solver in a list of all such sets.
2. During branching, the core solver checks if the current branching literal belongs to some such set S. If yes, the current branching literal is assigned true and all other theory case-split literals in S are assigned false. Otherwise, the default branching behaviour is used.
3. During propagation, the core solver may assign a truth value to a literal l in some set S of theory case-split literals. If so, the theory case-split check is invoked, i.e., the core solver checks whether two literals l_1, l_2 in the same set S have been assigned the value true. If this is the case, the core solver immediately generates the conflict clause $(\neg l_1 \vee \neg l_2)$.

4 Experimental Results

In this section, we describe the experimental evaluation of the Z3str3 solver to validate the effectiveness of the techniques presented in this paper. All techniques improve solver efficiency in isolation as well as in combination. In the interest of brevity we only report on the combined result in detail. We compare Z3str3 against four other state-of-the-art string solvers, namely, Z3str2 [20,19], CVC4 [8], S3 [16], and Norn [2], across industrial benchmarks obtained from Kaluza [12], PISA [15], and AppScan Source [1]. Each of these benchmark suites draw from real-world applications with diverse characteristics. All experiments were performed on a workstation running Ubuntu 15.10 with an Intel i7-3770k CPU and 16GB of memory. Also, we cross-verified the models generated by Z3str3 against Z3str2 and CVC4, and vice-versa.

Table 1 shows the summary of results for the Kaluza benchmark. Figure 1 presents the results in two cactus plots. Binaries for S3P are not publicly available so we were not able to evaluate it directly or include the timing information in the cactus plots in Figure 1. Instead, we report the aggregate results presented for
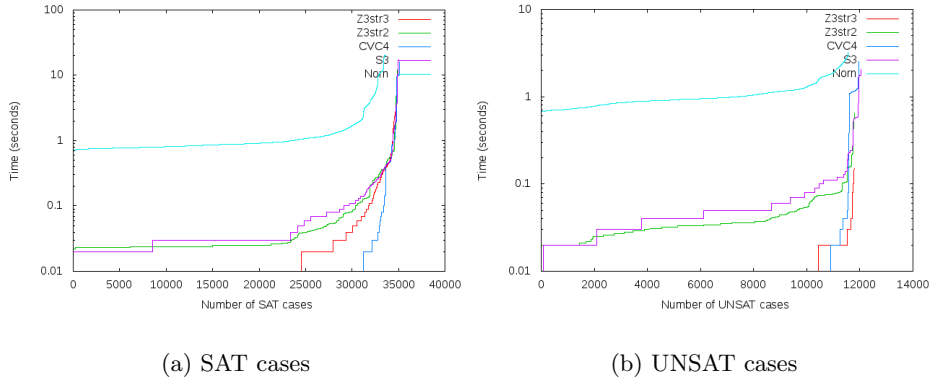
(a) SAT cases   (b) UNSAT cases

Fig. 1. Cactus plots for the Kaluza benchmark suite.

	Z3str3	Z3str2	CVC4	Norn	S3	S3P
sat	34885	34868	35128	33527	35016	35270
unsat	11786	11799	11957	11568	12049	12014
unknown	529	617	6	1913	0	0
timeout	84	0	0	276	219	0
error	0	0	193	0	0	0
crash	0	0	0	0	0	0
Total time (s)	5275.11	3997.63	4851.66	109280.76	10544.06	6972
Total time without timeouts (s)	3595.11	3997.63	4851.66	97784.00	6164.06	6972

Table 1. Results on cases from the Kaluza benchmark. Timeout=20 s.

this benchmark in the most recent S3P paper [17]. From Table 1 it is clear that Z3str3 is competitive with respect to CVC4, and is much faster than other tools. CVC4 has 193 errors, but Z3str3 times out on 84 cases. The unknowns in Z3str3 are because it lacks the feature to handle string equations with overlapping variables, similar to Z3str2.

Table 2 shows the results on the PISA benchmark. Norn was not able to solve any of the cases as it crashed upon seeing unrecognized string operators (e.g. `indexof`). From Table 2 we make the following observations. The tools Z3str3, Z3str2, and CVC4 are in agreement on all cases they are able to solve, with CVC4 and Z3str2 timing out on one SAT case which Z3str3 can solve in 16.58 seconds. The results for S3 are significantly worse; it is unable to solve `pisa-009.smt2` while the other three solvers all answer SAT reasonably quickly, and in addition S3 incorrectly answers UNSAT for `pisa-008.smt2`, `pisa-010.smt2`, and `pisa-011.smt2`, on which Z3str3 and (for two of these cases) Z3str2 and CVC4 all return SAT and produce a valid model.

Table 3 shows the results on the AppScan benchmark. Norn crashed on these cases as well upon seeing unrecognized string operators. From Table 3 we make the following observations. Z3str3, Z3str2, and CVC4 all agree on all cases they

input	Z3str3		Z3str2		CVC4		S3	
	result	time (s)	result	time (s)	result	time (s)	result	time (s)
pisa-000.smt2	sat	0.03	sat	0.25	sat	0.08	sat	0.07
pisa-001.smt2	sat	0.01	sat	0.19	sat	0.00	sat	0.07
pisa-002.smt2	sat	0.01	sat	0.10	sat	0.00	sat	0.05
pisa-003.smt2	unsat	0.00	unsat	0.02	unsat	0.01	unsat	0.02
pisa-004.smt2	unsat	0.01	unsat	0.05	unsat	0.39	unsat	0.05
pisa-005.smt2	sat	0.06	sat	0.14	sat	0.02	sat	0.04
pisa-006.smt2	unsat	0.01	unsat	0.05	unsat	0.32	unsat	0.05
pisa-007.smt2	unsat	0.01	unsat	0.05	unsat	0.37	unsat	0.05
pisa-008.smt2	sat	16.58	timeout	20.00	timeout	20.00	unsat X	4.73
pisa-009.smt2	sat	12.59	sat	0.62	sat	0.00	timeout	20.00
pisa-010.smt2	sat	0.03	sat	0.09	sat	0.00	unsat X	0.02
pisa-011.smt2	sat	0.04	sat	0.06	sat	0.00	unsat X	0.02

Table 2. PISA benchmark results. Timeout=20 s. X = incorrect response.

input	Z3str3		Z3str2		CVC4		S3	
	result	time (s)	result	time (s)	result	time (s)	result	time (s)
t01.smt2	sat	7.05	sat	1.31	sat	0.01	sat	0.23
t02.smt2	sat	0.13	sat	0.38	sat	0.01	unknown	0.04
t03.smt2	sat	0.53	sat	9.54	sat	3.82	sat X	0.14
t04.smt2	sat	0.68	sat	4.45	timeout	20.00	sat X	0.10
t05.smt2	sat	1.15	sat	16.84	sat	3.87	sat X	0.55
t06.smt2	sat	0.02	sat	0.15	sat	0.01	sat	0.13
t07.smt2	sat	2.62	sat	0.25	sat	0.00	unknown	0.02
t08.smt2	sat	0.01	sat	0.25	sat	0.17	sat X	0.03

Table 3. AppScan benchmark results. Timeout=20 s. X = incorrect response.

are able to solve. CVC4 performs better than Z3str3 on 3 cases and worse on 5 (including one timeout). Z3str2 performs better than Z3str3 on 2 cases and worse on 6, taking almost three times as long in total (33.17 seconds vs. 12.19 seconds). S3 returns UNKNOWN on two cases that are solved by the other three tools and produces invalid models which fail cross-validation for four other cases.

5 Discussion on Experimental Results, and Conclusions

The experimental results discussed here make clear the efficacy of theory-aware branching and case-split. The crucial insight behind these techniques is that biasing the search towards easier branches of the search tree (e.g., an arrangement that doesn't require splitting variables, as opposed to one with overlapping variables) is often very effective since most string constraints obtained from practical applications have the "small model" property. The slogan of theory-aware branching is "bias search towards easy cases first". We also note that Z3str3 and

CVC4 are sound, and more robust as compared to Norn and S3 which sometimes give wrong answers or crash on the benchmarks we used.

References

1. IBM Security AppScan Tool and Source. URL: http://www-03.ibm.com/software/products/en/appscan-source.
2. P. A. Abdulla, M. F. Atig, Y. Chen, L. Holík, A. Rezine, P. Rümmer, and J. Stenman. String constraints for verification. In *Proceedings of the 26th International Conference on Computer Aided Verification*, CAV'14, pages 150–166, 2014.
3. A. Aydin, L. Bang, and T. Bultan. *Automata-Based Model Counting for String Constraints*, pages 255–272. Springer International Publishing, Cham, 2015.
4. N. Bjørner, N. Tillmann, and A. Voronkov. Path feasibility analysis for string-manipulating programs. In *Proceedings of the 15th International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, TACAS '09, pages 307–321, 2009.
5. L. De Moura and N. Bjørner. Z3: an efficient smt solver. In *Proceedings of the Theory and practice of software, 14th international conference on Tools and algorithms for the construction and analysis of systems*, TACAS'08, pages 337–340, 2008.
6. M. Emmi, R. Majumdar, and K. Sen. Dynamic test input generation for database applications. In *ISSTA*, pages 151–162, 2007.
7. A. Kiezun, V. Ganesh, P. J. Guo, P. Hooimeijer, and M. D. Ernst. Hampi: A solver for string constraints. In *Proceedings of the Eighteenth International Symposium on Software Testing and Analysis*, ISSTA '09, pages 105–116, 2009.
8. T. Liang, A. Reynolds, C. Tinelli, C. Barrett, and M. Deters. A dpll(t) theory solver for a theory of strings and regular expressions. In *Proceedings of the 26th International Conference on Computer Aided Verification*, CAV'14, pages 646–662. 2014.
9. M. W. Moskewicz, C. F. Madigan, Y. Zhao, L. Zhang, and S. Malik. Chaff: Engineering an efficient SAT solver. In *Proceedings of the 38th Design Automation Conference, DAC 2001, Las Vegas, NV, USA, June 18-22, 2001*, pages 530–535. ACM, 2001.
10. G. Redelinghuys, W. Visser, and J. Geldenhuys. Symbolic execution of programs with strings. In *Proceedings of the South African Institute for Computer Scientists and Information Technologists Conference*, SAICSIT '12, pages 139–148, 2012.
11. P. Saxena, D. Akhawe, S. Hanna, F. Mao, S. McCamant, and D. Song. A symbolic execution framework for JavaScript. In *IEEE Symposium on Security and Privacy*, pages 513–528, 2010.
12. P. Saxena, D. Akhawe, S. Hanna, F. Mao, S. McCamant, and D. Song. A symbolic execution framework for javascript. In *Proceedings of the 2010 IEEE Symposium on Security and Privacy*, SP '10, pages 513–528, 2010.
13. R. Sebastiani. Lazy satisfiability modulo theories. volume 3, 2007.
14. K. Sen, S. Kalasapur, T. Brutch, and S. Gibbs. Jalangi: A selective record-replay and dynamic analysis framework for javascript. In *Proceedings of the 2013 9th Joint Meeting on Foundations of Software Engineering*, ESEC/FSE 2013, pages 488–498, New York, NY, USA, 2013. ACM.
15. T. Tateishi, M. Pistoia, and O. Tripp. Path- and index-sensitive string analysis based on monadic second-order logic. *ACM Trans. Softw. Eng. Methodol.*, 22(4):33:1–33:33, Oct. 2013.

16. M.-T. Trinh, D.-H. Chu, and J. Jaffar. S3: A symbolic string solver for vulnerability detection in web applications. In *Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security*, CCS '14, pages 1232–1243, 2014.

17. M.-T. Trinh, D.-H. Chu, and J. Jaffar. *Progressive Reasoning over Recursively-Defined Strings*, pages 218–240. Springer International Publishing, Cham, 2016.

18. G. Wassermann and Z. Su. Sound and precise analysis of web applications for injection vulnerabilities. In J. Ferrante and K. McKinley, editors, *PLDI*, pages 32–41. ACM, 2007.

19. Y. Zheng, V. Ganesh, S. Subramanian, O. Tripp, M. Berzish, J. Dolby, and X. Zhang. Z3str2: an efficient solver for strings, regular expressions, and length constraints. *Formal Methods in System Design*, pages 1–40, 2016.

20. Y. Zheng, V. Ganesh, S. Subramanian, O. Tripp, J. Dolby, and X. Zhang. *Effective Search-Space Pruning for Solvers of String Equations, Regular Expressions and Length Constraints*, pages 235–254. Springer International Publishing, Cham, 2015.

21. Y. Zheng, X. Zhang, and V. Ganesh. Z3-str: A z3-based string solver for web application analysis. In *Proceedings of the 2013 9th Joint Meeting on Foundations of Software Engineering*, ESEC/FSE 2013, pages 114–124, 2013.